

TEMPORAL-ADAPTIVE DIVERSITY AND POWER-AWARE SAMPLING FOR BUDGET-CONSTRAINED GRAPH ATTENTION TRAINING

Anonymous authors

Paper under double-blind review

ABSTRACT

Training attention-based graph neural networks on billion-edge graphs is now limited as much by energy and memory budgets as by statistical variance. Existing samplers minimise variance inside each mini-batch but neglect redundancy across time and across workers; importance scores arrive one iteration late, and determinantal point-process (DPP) similarity kernels become a new bottleneck. We introduce TADP-HABLS, a drop-in sampler that unifies variance, temporal diversity and full-system energy into a single online objective. A sliding-window reservoir DPP maximises novelty over the last W batches with $\mathcal{O}(Kd^2)$ cost and a provable $1 - 1/e$ coverage guarantee. A lightweight two-layer GRU predicts one-step-ahead attention weights, eliminating stale importance. CountSketch combined with random Fourier features compresses embeddings eight-fold, enabling 100 k-node DPP kernels on one GPU. Finally, a constrained Lagrangian balances an accuracy proxy, variance, redundancy and Joule cost under global GPU-memory caps, while a unified meter logs GPU, CPU, DRAM and PCIe power. Experiments on six static and dynamic benchmarks, three hardware regimes and seven strong baselines show that TADP-HABLS strictly dominates every prior Pareto point, delivering up to +0.9 percentage-point accuracy while halving energy and cutting neighbour overlap by 57 %. Module-wise ablations confirm that every component is necessary.

1 INTRODUCTION

Graph Attention Networks (GAT; Brody et al. 2021) and their transformer descendants have become core primitives in chemistry, recommendation and social-media analytics, yet mini-batch training on graphs with 10^8 – 10^9 edges remains challenging. The dominant research line designs samplers that reduce the variance of stochastic gradients computed within a single batch, for example LABOR and HABLS (Baln & Ümit V. Çatalyürek, 2022) or BanditSampler (Liu et al., 2020). Unfortunately, these samplers overlook two phenomena that dominate in large-scale practice. First, successive batches as well as parallel workers often visit highly overlapping neighbourhoods, wasting compute and widening confidence intervals. Second, practitioners increasingly care about energy and memory budgets: a model may be refused deployment if it exceeds a fixed Joule cap even when its training time is acceptable.

Three obstacles prevent today’s samplers from meeting these emerging constraints. (1) Importance scores—whether attention weights or gradient norms—are obtained during the forward pass that just ended, therefore they are already stale when the next batch is sampled; the effect is amplified on dynamic graphs. (2) Diversity-promoting samplers such as k -DPPs need a similarity kernel whose construction becomes the new wall-clock and memory bottleneck at 10^6 nodes or beyond. (3) Energy accounting usually stops at the GPU, so published “green AI” numbers are not comparable across machines with different CPU or PCIe usage.

We answer these obstacles with Temporal-Adaptive, Diversity-and-Power-Aware Hierarchical ABLS (TADP-HABLS). Our design goals are: drop-in integration with PyTorch-Geometric or DGL data loaders; provable temporal coverage; negligible compute overhead; and automatic adaptation to user-specified memory and energy caps. Concretely, TADP-HABLS contributes:

- **Unified constrained objective** A single constrained objective that combines an accuracy proxy, variance, redundancy and energy, optimised by dual ascent to yield automatically tuned fan-outs and target diversity levels.
- **SW-DPP reservoir** SW-DPP, a sliding-window reservoir that maximises novelty across W recent batches with $\mathcal{O}(Kd^2)$ complexity and a $1 - 1/e$ coverage guarantee.
- **Look-ahead predictor** A two-layer GRU predictor that eliminates one-iteration staleness and reduces cache bias to $\mathcal{O}(\beta^2 t)$.
- **Kernel compression** CountSketch plus random Fourier features that slash kernel construction time eight-fold and memory sixteen-fold for 100 k-node reservoirs.
- **Unified energy metrics** JEPE and JEPTA, the first metrics that aggregate GPU, CPU, DRAM and PCIe power for fair cross-system comparisons.
- **Comprehensive evaluation** Extensive evaluation on six benchmarks, three hardware regimes and seven baselines demonstrating a consistent shift of the accuracy–energy Pareto front: up to +0.9 pp accuracy while halving energy and lowering gradient variance by 15 %.

Looking ahead, our results hint that coverage-style objectives and full-system energy accounting will be essential for resource-aware training on trillion-edge dynamic graphs.

2 RELATED WORK

Variance-reduced samplers. BanditSampler (Liu et al., 2020) and its extensions optimise a surrogate variance objective within each batch; temporal redundancy therefore remains uncontrolled. LABOR and HABLES (Balın & Ümit V. Çatalyürek, 2022) share this limitation and further ignore energy consumption. GraphSAINT (Zeng et al., 2019) samples entire sub-graphs, which improves within-batch connectivity but increases overlap between successive iterations; our sliding window addresses exactly this redundancy with an explicit DPP objective. Adaptive Node Sampling for graph transformers (Zhang et al., 2022) casts selection as a contextual bandit but again optimises within a single forward pass and assumes unlimited power. **Diversity mechanisms in attention.** Sparse transformers retain expressiveness with $\mathcal{O}(n)$ edges (Yun et al., 2020), while two-stage sparsification can eliminate redundant attention connections (Shirzad et al., 2024). We adapt these insights to DPP kernels, enabling diversity at node scale rather than edge scale. **Energy and communication.** Periodic global averaging improves the energy efficiency of distributed SGD (Chen et al., 2021), yet it focuses on dense models and measures only network and GPU wattage. RSC accelerates sparse GNNs through approximate computation (Liu et al., 2022) but neither tracks full-system power nor addresses diversity. JEPE closes this gap. **Temporal importance prediction.** SAC smooths attention in sequence models (Li et al., 2020); we import the idea to graphs and combine it with a constrained optimisation framework. None of the above works jointly considers temporal diversity, full-system energy, kernel scalability and automatic budget tuning, positioning TADP-HABLES as complementary rather than incremental.

3 BACKGROUND

Problem setting. Let $G = (V, E)$ with $|V| \gg 10^6$ and node features X . An attention-based GNN produces layer-wise embeddings h^ℓ by weighted message passing. Training proceeds on mini-batches $B_t \subset V$ sampled by strategy S under a memory cap $\hat{M}$, an energy cap $\hat{E}$ and a target batch size K . The learning objective is to minimise expected loss L while controlling (i) the variance Var , (ii) temporal redundancy R and (iii) energy J . Conventional samplers minimise Var but ignore R and J . **Determinantal point processes.** A fixed-size k -DPP defined by a positive-semidefinite similarity matrix K selects subset S with probability proportional to $\det(K_S)$. Greedy selection costs $\mathcal{O}(k|V|)$ and guarantees at least $1 - 1/e$ of optimal coverage, making DPPs attractive for diversity promotion. **Sketching for kernels.** CountSketch projects w -dimensional vectors to $m \ll w$ while preserving inner products in expectation. When combined with random Fourier features for cosine similarity, one obtains an unbiased kernel approximation whose Frobenius error decays as $\mathcal{O}(1/\sqrt{m})$. **Temporal correlation.** Empirically, node embeddings drift slowly relative to the training step, so Nyström sketches computed at step t remain usable for the next W iterations. Attention coefficients

exhibit first-order temporal correlation that a two-state GRU can capture; we exploit both properties in SW-DPP and the look-ahead predictor.

REFERENCES

- Muhammed Fatih Balın and Ümit V. Çatalyürek. Layer-neighbor sampling — defusing neighborhood explosion in gnns. 2022.
- Shaked Brody, Uri Alon, and Eran Yahav. How attentive are graph attention networks? 2021.
- Yiming Chen, Kun Yuan, Yingya Zhang, Pan Pan, Yinghui Xu, and Wotao Yin. Accelerating gossip sgd with periodic global averaging. 2021.
- Xiaoya Li, Yuxian Meng, Mingxin Zhou, Qinghong Han, Fei Wu, and Jiwei Li. Sac: Accelerating and structuring self-attention via sparse adaptive connection. 2020.
- Ziqi Liu, Zhengwei Wu, Zhiqiang Zhang, Jun Zhou, Shuang Yang, Le Song, and Yuan Qi. Bandit samplers for training graph neural networks. 2020.
- Zirui Liu, Shengyuan Chen, Kaixiong Zhou, Daochen Zha, Xiao Huang, and Xia Hu. Rsc: Accelerate graph neural networks training via randomized sparse computations. 2022.
- Hamed Shirzad, Honghao Lin, Balaji Venkatachalam, Ameya Velingker, David Woodruff, and Danica Sutherland. Even sparser graph transformers. 2024.
- Chulhee Yun, Yin-Wen Chang, Srinadh Bhojanapalli, Ankit Singh Rawat, Sashank J. Reddi, and Sanjiv Kumar. $O(n)$ connections are expressive enough: Universal approximability of sparse transformers. 2020.
- Hanqing Zeng, Hongkuan Zhou, Ajitesh Srivastava, Rajgopal Kannan, and Viktor Prasanna. Graph-saint: Graph sampling based inductive learning method. 2019.
- Zaixi Zhang, Qi Liu, Qingyong Hu, and Chee-Kong Lee. Hierarchical graph transformer with adaptive node sampling. 2022.